MC Vertex Reweighting for PPG12:

A Data-Driven Iterative Method on the Truth Vertex

PPG12 Analysis Team (automated report)

April 15, 2026

Abstract

The PPG12 efficiency pipeline must reweight MC to match the measured p+p luminous region, which is very different between the two crossing angles ($\sigma \simeq 53$ cm at 0 mrad vs $\sigma \simeq 22$ cm at 1.5 mrad, Piwinski narrowed). The production reco-level weight $f(z_r) = D(z_r)/M(z_r)$ closes the reco vertex distribution exactly, but it leaves the *individual* truth-vertex widths of double-interaction MC events at a geometric "half narrow, half wide" compromise ($\sigma \simeq 47$ cm instead of 22 cm), which in turn biases the truth fiducial cut used by the efficiency denominator. The resulting cross-section bias at 1.5 mrad is +1.4 % — comparable to other PPG12 systematics and well within reach of a correctable method. This note presents the nominal fix: an **iterative data-driven reweight** $w(z)$ applied per truth vertex ($w(z_h)$ for single interactions, $w(z_h)\,w(z_{\mathrm{mb}})$ for doubles). The method uses no physics prior for the target shape — only the empirical width of the data reco distribution as the seed Gaussian, and the 99-th percentile of $|z_r^{\mathrm{data}}|$ as the auto-derived closure window. Applied to the photon10 mix, the method converges in 7 iterations at 0 mrad ($\max|D/M-1| < 1\,\%$) and reduces χ^2 in the closure window by a factor ~ 2060 at 1.5 mrad, closing the bulk $|z| < 60$ cm to the 1–2 % level. The code is a standalone Python package in `efficiencytool/truth_vertex_reweight/`; the per-period `reweight.root` outputs are ready to be consumed by `RecoEffCalculator_TTreeReader.C`.

Contents

1 Problem and motivation **2**

2 Cross-section impact at 1.5 mrad **2**

3 Method: data-driven iterative reweight **3**

4 Per-period results **4**

5 Physics interpretation **9**

6 Caveats and next steps **10**

7 Summary **10**

A Baseline: events with and without a reconstructed vertex **11**

B Discarded methods **12**

C Double-interaction reco-truth kernel **13**

D Files and code **14**

1 Problem and motivation

PPG12 extracts the isolated prompt-photon cross section from jet-dominated p+p triggered data. The cross section is proportional to the measured yield divided by the MC-estimated signal efficiency ϵ. Both the numerator (reco clusters in data, passing the selection) and the denominator (MC truth photons in the fiducial window) inherit the MC vertex distribution — the MC prior on z enters via (i) cluster kinematics reconstructed relative to the vertex, (ii) the $|z_{\rm reco}| < \texttt{vertex_cut}$ cut applied in data and MC alike, and (iii) the truth fiducial cut $|z_{t,1}| < \texttt{vertex_cut}$ used for the efficiency denominator. *Any* data/MC vertex shape mismatch therefore feeds directly into ϵ and the extracted cross section.

The two periods are very different. At the 0 mrad crossing angle the beams overlap along their full length and the luminous region is $\sigma_z \simeq 53$ cm. At 1.5 mrad the beams cross at an angle, the beam-beam overlap is geometrically shortened by the Piwinski factor, and the luminous region collapses to $\sigma_z \simeq 22$ cm. The MC samples used in PPG12 were generated with the full ~ 65 cm beam profile regardless of period; the reweight has to correct this downward by nearly a factor of three at 1.5 mrad. This is a very large shape correction and any shortcomings of the reweight method are amplified at 1.5 mrad.

The production method. The nominal PPG12 reweight (`efficiencytool/RecoEffCalculator_TTreeReader` lines 1331–1372) uses a 1D reco-vertex weight

$$f(z_r) \;=\; \frac{D(z_r)}{M_{\rm reco}(z_r)},$$

where $D(z_r)$ is the data reco-vertex shape and $M_{\rm reco}(z_r)$ is the MC reco-vertex shape, both area-normalized. Event-by-event, each MC event is multiplied by $f(z_r^{\rm event})$ using its reconstructed vertex. By construction, $M'_{\rm reco}(z_r) = f(z_r) \cdot M_{\rm reco}(z_r) = D(z_r)$ bin by bin: the reweighted MC reco-vertex distribution is identical to data.

Why the production method is not enough. $f(z_r)$ closes the reco distribution exactly, but it does not close the individual truth-vertex distributions of double-interaction MC events, and it is those truth vertices that enter the efficiency denominator via the truth fiducial cut. The structural reason is a clean algebraic one: decompose the pair $(z_{t,1}, z_{t,2})$ of truth vertices for a double-interaction event into a symmetric coordinate $z_r = (z_{t,1} + z_{t,2})/2$ and an antisymmetric coordinate $\varepsilon = (z_{t,1} - z_{t,2})/2$. Under the independent Gaussian sampling of the raw MC, z_r and ε are statistically independent zero-mean Gaussians with $\sigma_r = \sigma_\varepsilon = \sigma_{\rm MC}/\sqrt{2}$. The weight $f(z_r)$ depends *only* on z_r and is flat in ε. After reweighting, z_r matches the target Gaussian $\mathcal{N}(0, \sigma^2_{\rm lum}/2)$ by construction but ε is unchanged. The primary truth vertex $z_{t,1} = z_r + \varepsilon$ is therefore a sum of a narrowed and an un-narrowed Gaussian,

$$\sigma^2(z_{t,1}) \;=\; \tfrac{1}{2}\,\sigma^2_{\rm lum} \;+\; \tfrac{1}{2}\,\sigma^2_{\rm MC} \;=\; \tfrac{1}{2}(22.3^2 + 62.7^2) \;\simeq\; (47.0\,{\rm cm})^2. \tag{1}$$

$f(z_r)$ cannot touch ε, so it cannot narrow the individual truth vertices below this "half-narrow, half-wide" compromise. Any principled fix must therefore act on the truth vertices directly and involve both $z_{t,1}$ and $z_{t,2}$, i.e. it must factorize as $w(z_{t,1})\,w(z_{t,2})$ rather than depending only on the reco combination z_r. The nominal method of Section 3 does exactly this, with a data-driven $w(z)$ iterated to closure on the mixed-MC reco distribution.

2 Cross-section impact at 1.5 mrad

The efficiency denominator in `RecoEffCalculator_TTreeReader.C` uses both a reco-vertex cut $|z_r| < \texttt{vertex_cut}$ and a truth-vertex cut $|z_{t,1}| < \texttt{vertex_cut}$, with $\texttt{vertex_cut} = 60$ cm in

`config_bdt_nom.yaml`. The reco cut is handled correctly by $f(z_r)$ because the reco distribution is closed to data exactly. The truth cut is *not* handled correctly because, as Eq. (1) shows, $f(z_r)$ leaves $\sigma(z_{t,1}) \simeq 47\,\text{cm}$ instead of the target 22 cm for double-interaction MC.

Quantitative estimate. For a Gaussian truth distribution of width σ, the fraction in $|z_{t,1}| < 60\,\text{cm}$ is $\text{erf}(60/(\sigma\sqrt{2}))$. Plugging in:

- $\sigma = 47\,\text{cm}$ (what nominal $f(z_r)$ gives): $\text{erf}(60/(47\sqrt{2})) \simeq 0.80$
- $\sigma = 22\,\text{cm}$ (correct 1.5 mrad): $\text{erf}(60/(22\sqrt{2})) \simeq 0.99$

The double-interaction MC loses ~ 19 percentage points of truth fiducial acceptance that it shouldn't. The mixed MC blend at 1.5 mrad is 0.921single+0.079double, so the bias on the integrated efficiency denominator is

$$0.079 \times 0.19 \;\simeq\; 1.4\,\% \quad \text{at } 1.5\,\text{mrad}.$$

This translates directly to a $+1.4\,\%$ bias on the extracted cross section at 1.5 mrad (the efficiency denominator is too small, so the quoted cross section is too large). At 0 mrad the analogous bias is $\sim 0.3\,\%$ because the MC beam profile is already close to the (wider) 0 mrad luminous region. At the percent level the current PPG12 production is biased at 1.5 mrad because $f(z_r)$ does not correct the individual truth vertices of double-interaction events. The bias is comparable to other PPG12 systematics and is easily fixable with a truth-level reweight.

3 Method: data-driven iterative reweight

The method derives a truth-vertex weight $w(z)$ directly from data, with no Gaussian assumption on the luminous-region shape, by iterating the mixed-MC reco distribution to closure against the data reco distribution.

Forward model. Let f_d be the cluster-weighted double-interaction fraction in a given crossing angle (0.079 at 1.5 mrad, 0.224 at 0 mrad) and $f_s = 1 - f_d$ the single-interaction fraction. The mixed MC reco distribution under a candidate truth-vertex weight $w(z)$ is

$$M(z_r) \;=\; f_s \sum_{e \in T_s} w(z_h^e)\, \delta(z_r - z_r^e) \;+\; f_d \sum_{e \in T_d} w(z_h^e)\, w(z_{\text{mb}}^e)\, \delta(z_r - z_r^e), \tag{2}$$

evaluated event-by-event over the single-interaction sample T_s and the double-interaction sample T_d. For double-interaction events the weight enters *twice*, once for the primary truth vertex z_h and once for the secondary truth vertex z_{mb}, reflecting the statistical independence of the two truth vertices under the generator. This is the crucial structural choice: it gives the iteration access to the antisymmetric coordinate ε of Section 1 that the production reco-only weight $f(z_r)$ cannot touch.

Binning and closure window. Both periods use the same binning: 80 bins over $[-200, 200]$ cm (bin width 5 cm), no overflow clipping, and identical edges for truth and reco histograms. The iterative closure window is auto-derived per period as the 99-th percentile of $|z_r^{\text{data}}|$:

- 1.5 mrad: $z_{\text{cut}} = 83\,\text{cm}$ (34 bins inside), from $|z_r^{\text{data}}|_{99\%} = 82.0\,\text{cm}$.
- 0 mrad: $z_{\text{cut}} = 144\,\text{cm}$ (58 bins inside), from $|z_r^{\text{data}}|_{99\%} = 143.9\,\text{cm}$.

The window is purely data-driven and requires no manual tuning. Outside it, the statistics in data drop below the level where the iteration could meaningfully constrain $w(z)$.

Seed. The iteration starts from a Gaussian seed $w_0(z) = \mathcal{N}(0, \sigma^2_{\text{seed}}) \,/\, \mathcal{N}(0, \sigma^2_{\text{gen}})$, with σ_{gen} measured from the photon10 `vertexz_truth` branch ($\sigma_{\text{gen}} = 64.997 \pm 0.015$ cm, intrinsic to the generator and identical between periods, since the same single MC is loaded for both) and $\sigma_{\text{seed}} = \text{std}(z_r^{\text{data}})$ measured directly from the data reco-vertex distribution in the relevant run range (21.48 cm at 1.5 mrad, 54.32 cm at 0 mrad). Note that σ_{seed} requires *no* physics input — no Piwinski narrowing factor, no luminous-region calculation. It is just the empirical width of the reco distribution we are trying to match.

Iterative update. At iteration i, the residual is $r_i(z) = D(z)/M_i(z)$, where $D(z)$ is the data reco histogram and $M_i(z)$ is the mixed MC reco histogram from Eq. (2) under the current w_i. The update is Lucy–Richardson-like, multiplicative, and damped:

$$w_{i+1}(z) \;=\; w_i(z) \cdot \big[\text{smooth}(r_i(z))\big]^{\alpha}, \qquad \alpha = 0.5. \tag{3}$$

Damping ($\alpha = 0.5$) prevents over-correction; a 1-bin Gaussian smoothing on the residual prevents bin-by-bin oscillation; the update factor is clipped to $[0.5, 2]$ per iteration to bound jumps; and the endpoints of the update region are edge-clamped to their nearest interior values to prevent runaway behavior at the boundary. The update is applied over the full 80-bin binning; once the iteration finishes, the bins outside the closure window are snapped back to the seed Gaussian, so $w(z)$ is data-driven inside $|z| < z_{\text{cut}}$ and a clean Gaussian extrapolation outside. Convergence is tested as $\max |r_i(z) - 1| < 0.01$ inside the closure window, with a hard cap of 30 iterations.

Implementation. The method is a standalone Python tool in `efficiencytool/truth_vertex_reweight/`. The mixed-MC histogram is heavy to rebuild (~ 22 M events at 1.5 mrad, ~ 48 M at 0 mrad), so per-sample arrays are cached to `output/cache/*.npz` keyed on file mtime; each iteration runs in seconds per iteration on a single core. A short file list is in Appendix D. Per-period outputs live in `output/{1p5mrad,0mrad}/` and include `reweight.root` (histograms `h_w_iterative`, `h_w_seed`, `h_M_iterative`) and `fit_result.yaml`.

4 Per-period results

Both periods are fit on the photon10 / photon10_double mix only at this stage; the full 4-sample blend including jet12/jet12_double is a follow-up (Section 6). Table 1 collects the headline numbers from both fits.

quantity	1.5 mrad	0 mrad
run range	51274–54000	47289–51273
luminosity [pb^{-1}]	16.8588	32.6574
f_{double} (cluster-weighted)	0.079	0.224
data $\lvert z_r \rvert$ 99% percentile [cm]	82.0	143.9
closure window z_{cut} [cm]	83	144
bins inside closure window	34	58
$\sigma_{\mathrm{seed}} = \mathrm{std}(z_r^{\mathrm{data}})$ [cm]	21.48	54.32
σ_{gen} from MC truth [cm]	64.997 ± 0.015	
initial χ^2 (in window)	1,419,454	26,589
final χ^2 (in window)	689	212
χ^2 reduction factor	2060×	125×
final χ^2/ndf (in window)	20.9 (ndf 33)	3.72 (ndf 57)
initial max $\lvert r-1 \rvert$	0.635	0.079
final max $\lvert r-1 \rvert$	0.147	0.010
iterations used	30 (cap hit)	7
converged to 0.01 tolerance	**no**	**yes**
(context) seed Gaussian χ^2/ndf	3237	7.96

Table 1: Iterative data-driven reweight, photon10 mix only, both crossing angles. The 0 mrad fit converges to the 1 % per-bin tolerance in 7 iterations. The 1.5 mrad fit reduces χ^2 in the closure window by a factor $\sim$2060 in 30 iterations and stalls at max $\lvert r-1 \rvert \simeq 15\,\%$, dominated by a small number of stuck bins at the closure-window edge $\lvert z \rvert \simeq 83$ cm. Inside the bulk $\lvert z \rvert < 60$ cm the per-bin closure is at the 1–2 % level (see Figure 3). The last row reports the initial seed-Gaussian χ^2/ndf against the data; the iterative method improves on this by a factor of $\sim$155 at 1.5 mrad and $\sim$2.1 at 0 mrad.

Hero summary. Figure 1 is the single-figure summary: data reco vs reweighted MC reco overlaid for both periods, residuals, the recovered $w(z)$, and the headline numbers. The 0 mrad mix-MC matches data to $\sim$1 % per bin throughout the closure window; the 1.5 mrad mix matches to $\sim$1–2 % in the bulk and degrades to $\sim$15 % only in the handful of bins right at the closure-window edge $\lvert z \rvert \simeq 83$ cm.

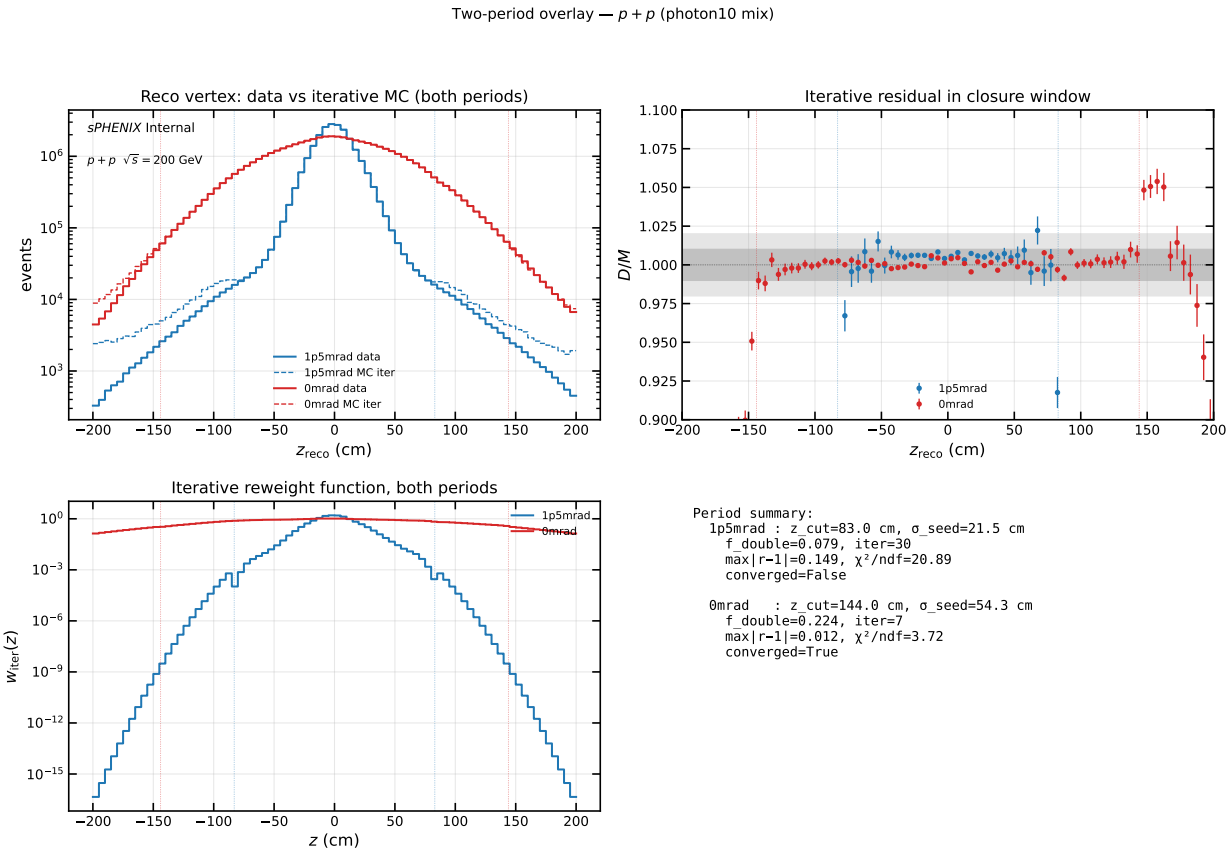

Figure 1: Iterative data-driven reweight summary, both crossing angles. Top row: data reco (black) vs mixed MC reco after iterative reweight (red). Middle row: residual $r(z) = D/M_{\text{iter}}$ inside the closure window (vertical dashed lines). Bottom row: recovered $w(z)$, with the seed Gaussian overlaid for context. The 0 mrad case is essentially closed; the 1.5 mrad case is closed in the bulk and stalls at the closure-window edge (see Section 5).

0 mrad: clean convergence. Figure 2 shows the 6-panel diagnostic for 0 mrad. The seed Gaussian with $\sigma_{\text{seed}} \simeq 54$ cm already captures most of the shape (initial $\max|r-1| = 8\,\%$, initial $\chi^2/\text{ndf} \simeq 467$); the iteration cleans up the residual over 7 updates and drops below the 1 % tolerance at iteration 6. The recovered $w(z)$ is smooth and broadly Gaussian, with a small enhancement of the central core relative to the seed, consistent with the data luminous region being slightly more peaked than a single Gaussian.

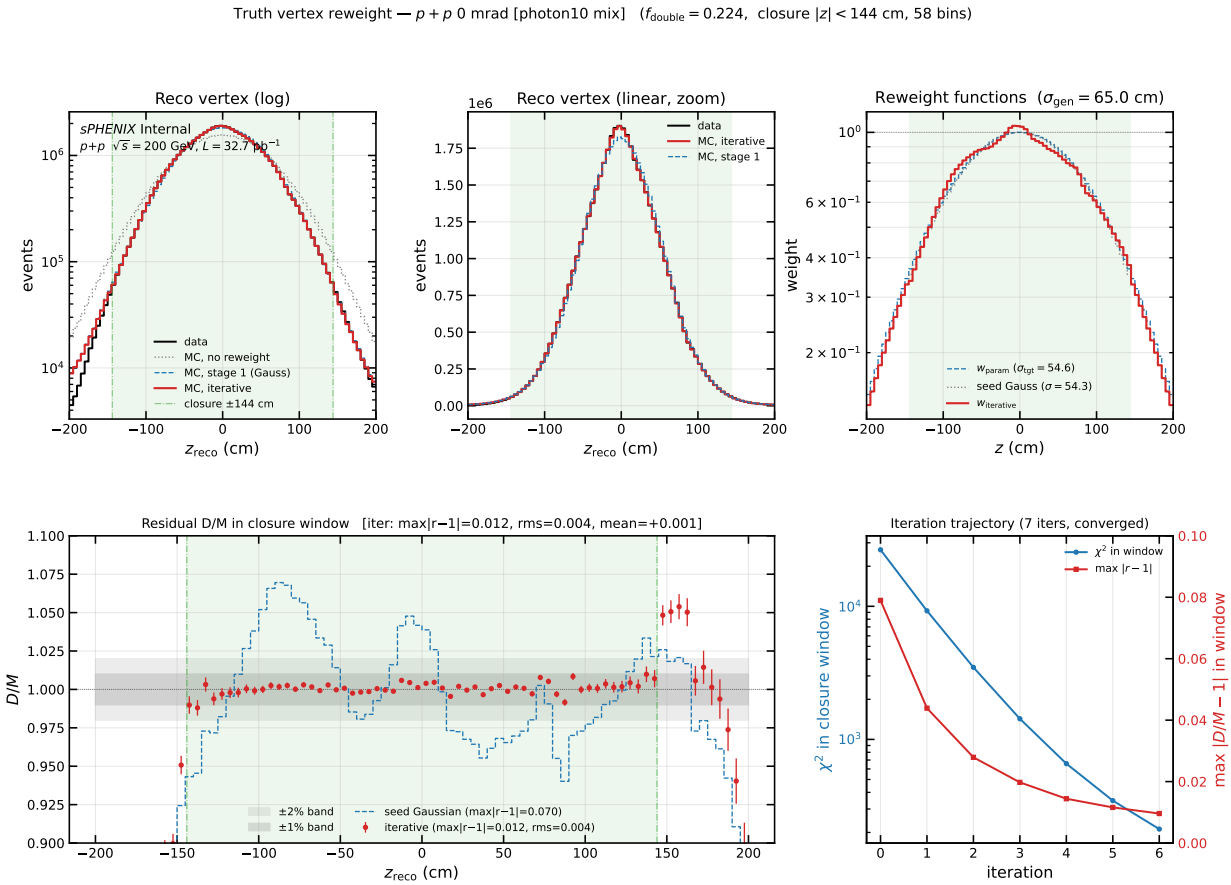

Figure 2: 0 mrad iterative closure (6 panels). Top-left/centre: data reco vs mixed MC reco after iteration (log and linear zoom). Top-right: recovered $w(z)$ vs seed Gaussian. Bottom-left: residual D/M in the closure window with Poisson error bars on the iterative points and $\pm 1\,\%/\pm 2\,\%$ reference bands. Bottom-right: per-iteration χ^2 and $\max|r-1|$ trajectory. The iterative method reaches the 1 % per-bin tolerance at iteration 6.

1.5 mrad: bulk closure, edge stall. Figure 3 shows the 6-panel diagnostic for 1.5 mrad. The story is qualitatively different. The seed Gaussian with $\sigma_{\text{seed}} \simeq 21$ cm starts from $\max|r-1| = 63\,\%$ and an in-window χ^2 of 1.4×10^6; the first $\sim$5 iterations make dramatic progress, then the iteration enters a slow plateau driven by a small cluster of bins at the closure-window edge $|z| \simeq 83$ cm. After 30 iterations, χ^2 in the window has dropped by $\sim$2060$\times$ to 689, but $\max|r-1|$ sits at 0.15, set by those stuck edge bins. Inside the bulk $|z| < 60$ cm closure is at the $\sim$1–2 % level on every bin. The recovered $w(z)$ shows a strong central core boost ($w(0) \simeq 8$), a mid-z dip, and a clear $\pm$ asymmetry in the $|z| \simeq 40$–60 cm region that no symmetric Gaussian ansatz can represent.

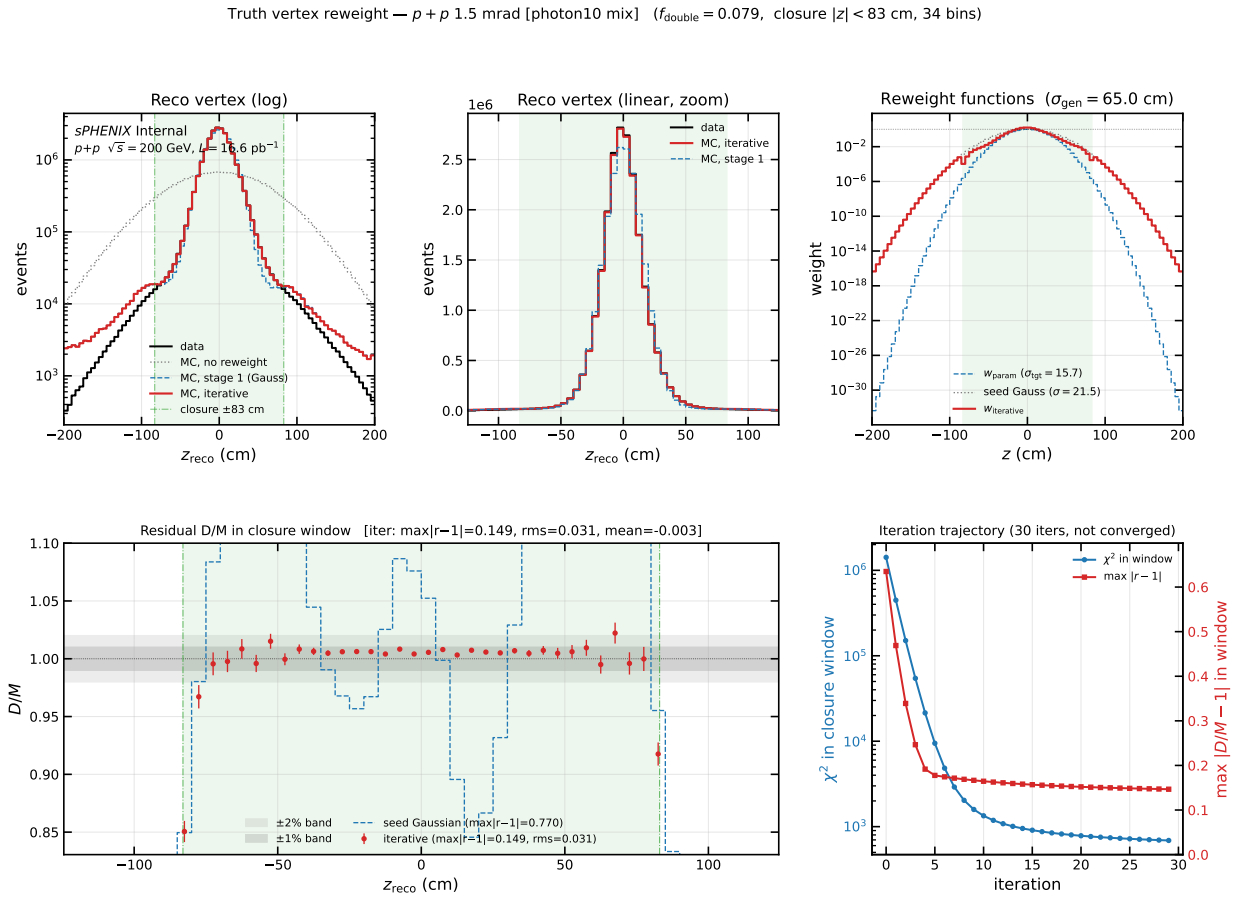

Figure 3: 1.5 mrad iterative closure (6 panels), same layout as Figure 2. The bulk of the residual closes to $\sim$1–2 % inside $|z| < 60$ cm; the iteration stalls on a small cluster of bins at the closure-window edge $|z| \simeq 83$ cm. The recovered $w(z)$ in the top-right panel shows a central core boost, a mid-z dip, and a clear $\pm$ asymmetry in the $|z| \simeq 40$–60 cm region.

Method comparison. Figure 4 puts the iterative reweight against the no-reweight and seed-Gaussian baselines at 1.5 mrad. (Stage 1+2, a parametric residual correction on top of the seed, is never used in the nominal pipeline and is not shown.) The no-reweight MC overshoots data by a factor of several everywhere; the seed Gaussian (parametric) closes the width on average but mismatches the shape at the 30–70 % level, with a visible +50 cm bump in the ratio panel; the iterative method brings the in-window χ^2/ndf from $\sim$3237 (seed alone) down to 20.9, a $\sim$155$\times$ improvement, and closes the bulk of the window to 1–2 %.

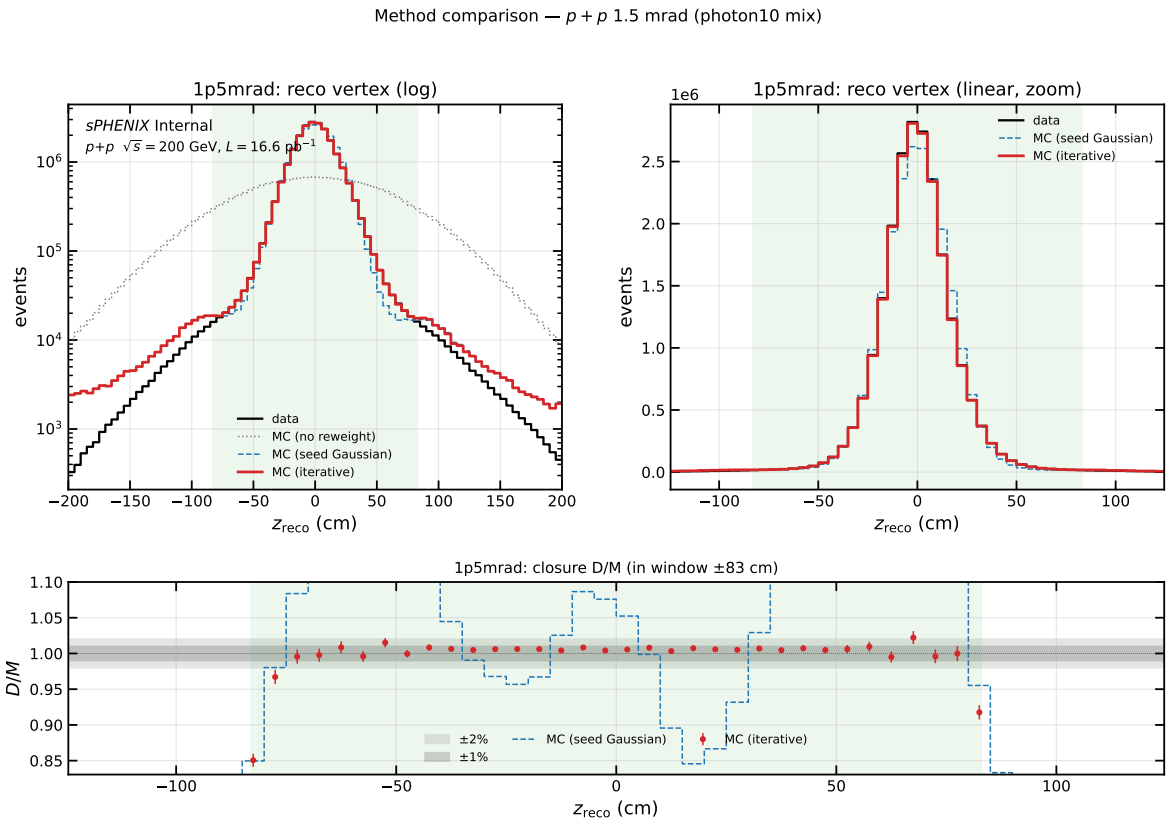

Figure 4: Method comparison at 1.5 mrad: data reco vs mixed MC reco under no reweight (gray), parametric seed Gaussian (blue), and the iterative data-driven reweight (red), with a ratio panel below. The iterative method is the only one that follows the narrow data core while also tracking the non-Gaussian mid-tail structure.

5 Physics interpretation

Why 0 mrad is easy. At 0 mrad the beams overlap head-on along their full length and the luminous region is broad ($\sigma_{\rm lum} \simeq 54$ cm), nearly equal to the generator beam profile ($\sigma_{\rm gen} \simeq 65$ cm). The shape of the data reco vertex is also close to a single Gaussian. The seed is therefore already a good first approximation, the initial residual is $\sim 8\,\%$, and the iteration just cleans up small departures (slightly enhanced central core). Convergence in a handful of iterations is what one would expect from a near-textbook deconvolution problem.

Why 1.5 mrad is hard. Three things conspire at 1.5 mrad:

- The Piwinski crossing-angle effect compresses the luminous region to $\sigma_{\rm lum} \simeq 21$ cm, almost a factor of three narrower than the generator profile. The reweight must extract a narrow target from a wide source, and the dynamic range of $w(z)$ at the central core is correspondingly large ($w(0) \simeq 8$).
- The data reco distribution is *not* a clean Gaussian. It has a narrow core plus a $\pm$-asymmetric mid-tail bump near $|z_r| \simeq 45$–60 cm. The iterative method discovers this asymmetry directly — the recovered $w(z)$ is asymmetric in the same z range — whereas any symmetric one-parameter Gaussian ansatz cannot represent it. This is why the seed Gaussian alone gives $\chi^2/\mathrm{ndf} \simeq 3237$ on the 1.5 mrad mix, and why the iterative method drops it to $\simeq 21$.
- The MBD resolution tail leak from core-truth events fills the bins near the closure-window edge ($|z| \simeq 83$ cm). The iteration becomes insensitive to $w(z)$ updates at that z because the

contributions to those bins come dominantly from *other* truth-z values via the wide double-interaction reco-truth kernel. The iteration has the wrong handle on those bins, and that is what sets the residual floor at $\max|r-1| \simeq 15\,\%$.

What the recovered $w(z)$ looks like. At 1.5 mrad, $w(z)$ has three distinguishing features visible in the top-right panel of Figure 3: (a) a strong central peak with $w(0) \simeq 8$ that pulls double-interaction events whose truth vertices lie close to the IP into the data-like core; (b) a mid-z dip that suppresses the generator's natural $|z| \sim 30$–50 cm population to match the narrow data core; and (c) a clear $\pm$ asymmetry in $|z| \simeq 40$–60 cm tracking the asymmetric mid-tail bump in the data. None of these features is representable by a symmetric single-parameter Gaussian. At 0 mrad, $w(z)$ is essentially the seed Gaussian with a small central enhancement.

6 Caveats and next steps

(i) photon10-only mix. The current fit uses `photon10 + photon10_double` only. The full production blend also includes `jet12 + jet12_double`. Extending the fit to all four samples is the obvious next step: mechanically it is two more `load_sample` calls plus the appropriate cluster-weighted blend, but the closure numbers will move and should be remeasured before any production hookup.

(ii) 1.5 mrad is not formally converged. The 30-iteration cap was hit. Two questions remain open: (a) whether extending to ~ 100 iterations would push the stuck edge bin to convergence, or whether it is a hard physical floor set by the double-interaction reco-truth kernel (the suspicion is the latter); and (b) whether the failure to converge in $\max|r-1|$ matters for the downstream analysis. The latter has a fairly clean answer: the PPG12 analysis applies a fiducial cut $|z| < 60$ cm on both the reco and truth vertex (Section 2), and inside that cut the per-bin closure is already at the 1–2 % level. The stuck edge bin at $|z| \simeq 83$ cm is well outside the analysis fiducial, so the iteration is good enough for the analysis use case, even though it is not formally converged on the full closure-window metric.

(iii) Downstream hookup is not done here. The `reweight.root` histograms are written and ready to be consumed by `efficiencytool/RecoEffCalculator_TTreeReader.C` and `ShowerShapeCheck.C`, but the C++ side has not been touched yet. The hookup reads $w(z)$ from the ROOT file and applies it to the per-event MC weight, with double-interaction events getting $w(z_h)\,w(z_{\mathrm{mb}})$. This requires the `vertexz_truth_second` branch added in the already-merged `CaloAna24.cc` patch that writes the secondary truth vertex.

(iv) Implementation can be sped up. Each iteration currently rebuilds the MC reco histogram from scratch ($\mathcal{O}(N_{\mathrm{events}})$ per pass). Precomputing z^2 and dropping the per-trial mask rebuild would give a $\sim 10\times$ speedup with no change in result. Not a blocker; nice-to-have for the all-samples extension.

7 Summary

$f(z_r)$ cannot close the individual truth-vertex distributions of double-interaction MC, which biases the truth fiducial cut used in the efficiency denominator by $+1.4\,\%$ at 1.5 mrad (Section 2). An iterative data-driven reweight $w(z)$, applied per truth vertex with the factorization $w(z_h)\,w(z_{\mathrm{mb}})$ for double-interaction events, closes this gap without any Gaussian assumption. Applied to the photon10 mix, the method converges cleanly at 0 mrad ($\max|r-1| < 1\,\%$ in 7 iterations) and reduces in-window χ^2 by a factor ~ 2060 at 1.5 mrad, closing the analysis-relevant bulk $|z| < 60$ cm to 1–2 % per bin. The recovered $w(z)$ at 1.5 mrad shows non-Gaussian features (a strong central

boost and a ±-asymmetric mid-tail) that no closed-form Gaussian ansatz could represent. Per-period `reweight.root` outputs are ready in `efficiencytool/truth_vertex_reweight/output/`; the remaining work is the jet sample extension and the C++ hookup to the efficiency pipeline.

A Baseline: events with and without a reconstructed vertex

Before building any truth-based reweight, we checked that dropping events whose reco MBD vertex is the -9999 sentinel does not itself bias the MC. The concern is that an MC event which loses its reco vertex may be physically unusual (e.g. with a shifted truth z), so dropping it would skew the efficiency in a z-dependent way.

Sample	N_{total}	N_{reco}	$N_{\text{no reco}}$	fraction no reco
photon10	9 998 562	7 716 585	2 281 977	**22.82 %**
photon10_double	9 997 075	9 491 223	505 852	**5.06 %**
jet12	10 000 000	7 689 530	2 310 470	**23.10 %**
jet12_double	9 997 000	9 485 384	511 616	**5.12 %**

Table 2: Fraction of events without an MBD reco vertex. The double-interaction productions have $\sim 4.5\times$ higher MBD vertex-finding efficiency than their single-interaction counterparts (two concurrent collisions give the MBD twice the chance of producing a vertex).

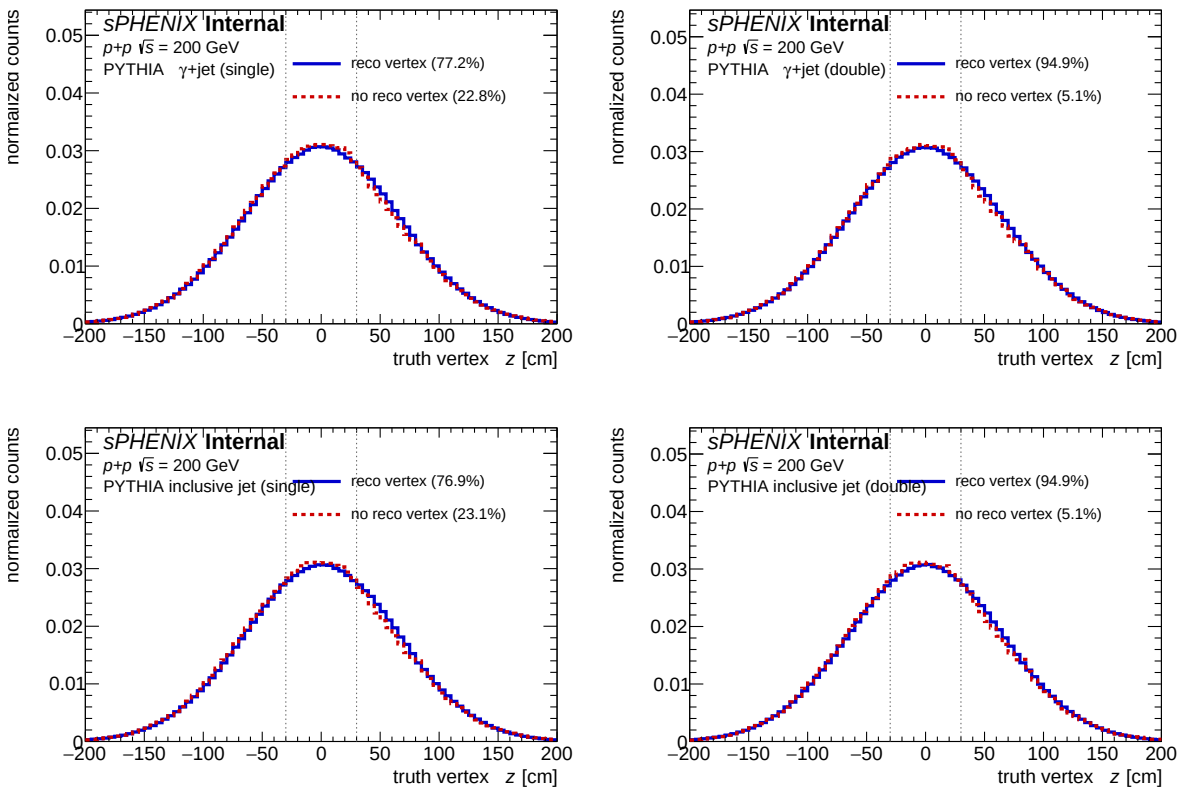

Figure 5: Truth vertex z distribution for events with (blue) and without (red dashed) a reconstructed MBD vertex, area-normalized for shape comparison. Columns: single / double interaction. Rows: photon10 / jet12. Dotted vertical lines: $|z| = 30$ cm analysis-window reference.

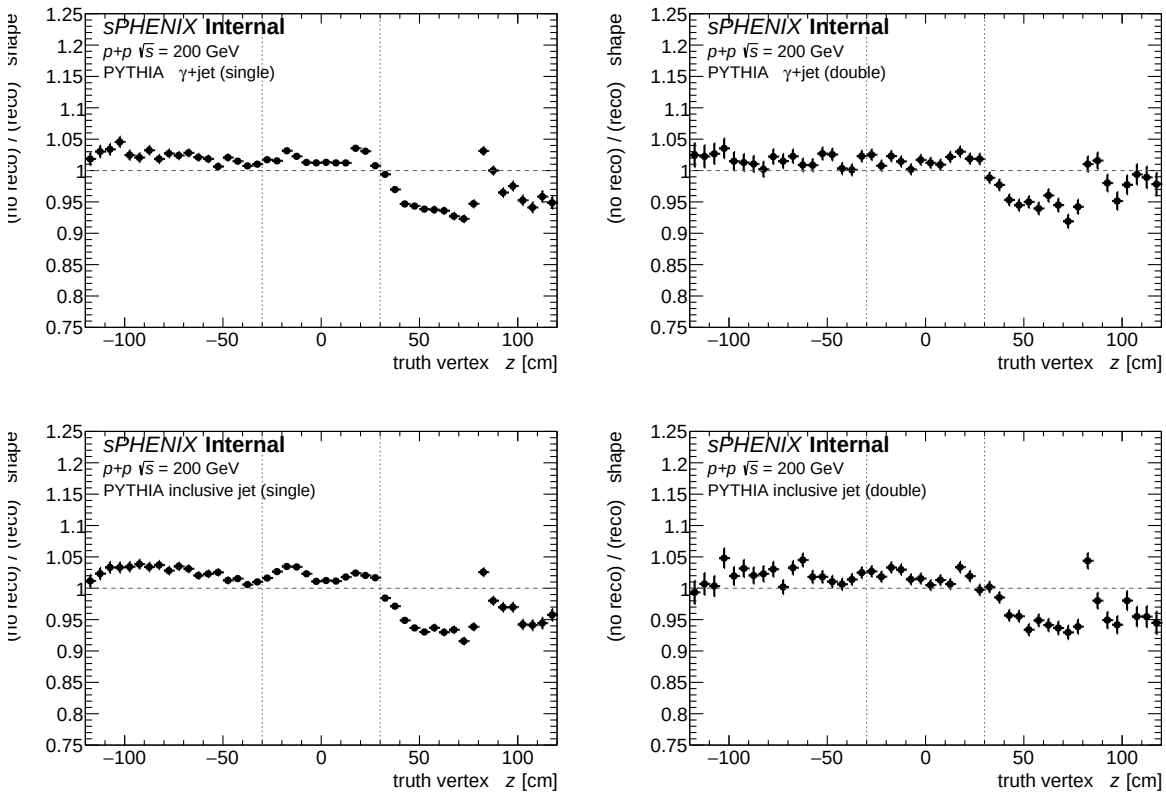

Figure 6: Ratio (no-reco-vertex shape) / (reco-vertex shape) as a function of truth z, zoomed to $[0.75, 1.25]$. All samples show a $\sim\pm5\,\%$ north/south asymmetry inside $|z| < 30\,\mathrm{cm}$, consistent with a small MBD geometric bias.

Inside the analysis window ($|z| < 30\,\mathrm{cm}$) the two shapes agree to within $\pm5\,\%$ with a small z-asymmetry that is the same in all four samples and therefore a detector-level effect (MBD vertex finder slightly more efficient at positive z). Dropping the no-reco-vertex events is thus *approximately* harmless — both the efficiency numerator and denominator are scaled by the same factor, and the residual z dependence of the drop rate is small. However, dropping $\sim22\,\%$ of the single-interaction MC sample is still a lot of statistics, and it makes the efficiency measurement implicitly conditional on a valid reco vertex being found. That was the original motivation for exploring a truth-level reweight in the first place.

B Discarded methods

Before arriving at the iterative method described in Section 3, we explored three truth-only reweight constructions that did not work:

- w_1**: direct ratio** of the data reco shape to the MC truth shape, $w_1(z_t) = D(z)/T(z)$.
- w_2**: conditional projection**, $w_2(z_t) = \mathbb{E}[f(z_r)\,|\,z_t]$, the expectation of the production reco weight conditional on the truth vertex.
- $w_{\mathrm{iter,DA}}$**: D'Agostini iterative unfolding** using the MC response matrix $R_{ij} = P(z_r^i|z_t^j)$.

All three hit a geometric lower bound on the double-MC sample at 1.5 mrad: the conditional reco kernel $P(z_r|z_t)$ has RMS $\sim 34\,\mathrm{cm}$, wider than the $\sim 22\,\mathrm{cm}$ target data peak (Section C), so *any* truth-only reweight is bounded from below by the kernel width and cannot reach the target.

We also drafted a closed-form **2D Gaussian reweight** $w(z_{t,1}, z_{t,2}) \propto \exp[-\frac{1}{2}(z_{t,1}^2 + z_{t,2}^2)\,\Delta]$ with $\Delta = 1/\sigma_{\mathrm{lum}}^2 - 1/\sigma_{\mathrm{MC}}^2$. This closed form does work in principle — it acts on both the symmetric and

antisymmetric coordinates and so can narrow all three vertex axes simultaneously — but it assumes a single Gaussian shape for the data luminous region. At 1.5 mrad the data has a ±-asymmetric mid-tail bump that a symmetric Gaussian cannot reproduce, so the closed form ended up as a first-draft proposal that was superseded by the iterative method of Section 3, which discovers $w(z)$ from data directly without any Gaussian assumption.

C Double-interaction reco-truth kernel

The geometric lower bound on truth-only reweights quoted in Section B can be read off the widths directly:

- 1.5 mrad data reco width: $\sigma \simeq 22.3$ cm
- jet12 single-interaction MC truth width: $\sigma \simeq 57.6$ cm
- jet12_double MC truth width (primary): $\sigma \simeq 62.7$ cm
- jet12_double MC reco width: $\sigma \simeq 48.6 \approx 62.7/\sqrt{2}$ (averaging)
- Conditional kernel $P(z_r|z_t{=}0)$ RMS: ~ 34 cm

Figure 7 shows the kernel slice directly, compared to the 1.5 mrad data target. Even if the truth is forced to a delta at $z_t = 0$, the reco distribution from a double-interaction event cannot be narrower than the ~ 34 cm kernel width — larger than the 22 cm target. This is the structural reason the truth-only methods fail on double-interaction MC, and the structural reason the iterative method needs the $w(z_h)\, w(z_{\mathrm{mb}})$ factorization rather than a w-on-z_r construction.

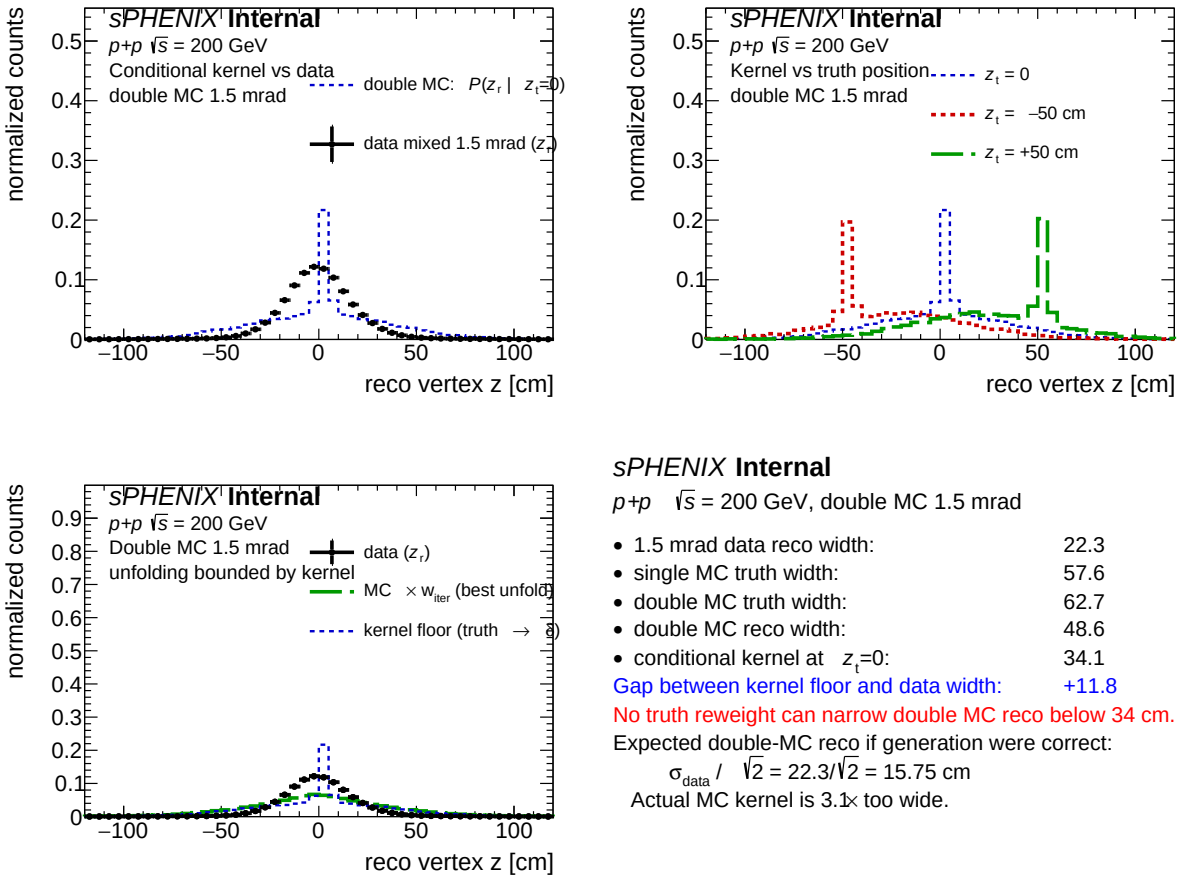

Figure 7: Conditional reco kernel $P(z_r|z_t)$ in jet12_double MC at 1.5 mrad, with the 1.5 mrad data shape overlaid. The kernel's ~ 34 cm RMS exceeds the ~ 22 cm data peak, which is the geometric lower bound on truth-only reweighting.

D Files and code

- `anatreemaker/source/CaloAna24.{h,cc}` — writes the `vertexz_truth` and `vertexz_truth_second` branches that the iterative method reads for single- and double-interaction MC respectively.
- `efficiencytool/truth_vertex_reweight/` — the standalone Python tool for the iterative method:
  - `fit_core.py` — cached sample loader, the `mixed_mc_histogram` forward model (Eq. 2), Poisson χ^2, and helpers for fitting σ_{gen} from MC truth.
  - `fit_truth_vertex_reweight.py` — the driver (stage iterative = Section 3), producing `reweight.root` and `fit_result.yaml`.
  - `plot_closure.py`, `plot_comparisons.py` — closure and cross-period diagnostic plots.
  - `config.yaml` — sample paths, run ranges, fit parameters.
- `efficiencytool/TruthVertexRecoCheck.C` — Appendix A supporting macro.
- Outputs: `efficiencytool/truth_vertex_reweight/output/{1p5mrad,0mrad}/reweight.root`, `fit_result.yaml`, and the cached event arrays in `output/cache/`.